

QUESTION 1

MapReduce Question :-

*****code*****

```
import java.io.*;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.mapreduce.lib.input.*;
import org.apache.hadoop.mapreduce.lib.output.*;

public class AllTimeHigh {
    public static class MapClass extends Mapper<LongWritable,Text,Text,DoubleWritable>
    {
        private Text stock_id = new Text();
        private DoubleWritable High = new DoubleWritable();

        public void map(LongWritable key, Text value, Context context)
        {
            try{
                String[] str = value.toString().split(",");
                double high = Double.parseDouble(str[4]);
                stock_id.set(str[1]);
                High.set(high);

                //context.write(new Text(stock_id),new LongWritable(vol));
                context.write(stock_id, High);
            }
            catch(Exception e)
            {
                System.out.println(e.getMessage());
            }
        }
    }
}
```

```

    public static class ReduceClass extends Reducer<Text,DoubleWritable,Text,DoubleWritable>
    {
        private DoubleWritable result = new DoubleWritable();

        public void reduce(Text key, Iterable<DoubleWritable> values,Context context) throws
IOException, InterruptedException {
            double maxVal=0;
            double temp_val=0;

            for (DoubleWritable value : values) {
                temp_val = value.get();
                if (temp_val > maxVal) {
                    maxVal = temp_val;
                }
            }
            result.set(maxVal);

            context.write(key, result);
            //context.write(key, new LongWritable(sum));
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        //conf.set("name", "value")
        //conf.set("mapreduce.input.fileinputformat.split.minsize", "134217728");
        Job job = Job.getInstance(conf, "Highest Price for each stock");
        job.setJarByClass(AllTimeHigh.class);
        job.setMapperClass(MapClass.class);
        //job.setCombinerClass(ReduceClass.class);
        job.setReducerClass(ReduceClass.class);
        job.setNumReduceTasks(1);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(DoubleWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

```
select count(*) from nyse;
```

```
2022-06-20 09:39:42,945 Stage-1 map = 0%, reduce = 0%
```

```
2022-06-20 09:39:57,133 Stage-1 map = 100%, reduce = 0%, Cumulative
CPU 4.75 sec
```

```
2022-06-20 09:40:06,416 Stage-1 map = 100%,  reduce = 100%,
Cumulative CPU 7.56 sec
MapReduce Total cumulative CPU time: 7 seconds 560 msec
Ended Job = job_1654490426372_5707
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1    Cumulative CPU: 7.56 sec    HDFS
Read: 41000131 HDFS Write: 106 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 560 msec
OK
735026
Time taken: 59.27 seconds, Fetched: 1 row(s)
```

```
select stock_id, max(high) from nyse group by stock_id;
```

```
Hadoop job information for Stage-1: number of mappers: 1; number of
reducers: 1
2022-06-20 09:44:09,105 Stage-1 map = 0%,  reduce = 0%
2022-06-20 09:44:58,510 Stage-1 map = 100%,  reduce = 0%, Cumulative
CPU 4.98 sec
2022-06-20 09:45:08,917 Stage-1 map = 100%,  reduce = 100%,
Cumulative CPU 8.73 sec
MapReduce Total cumulative CPU time: 8 seconds 730 msec
Ended Job = job_1654490426372_5731
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1    Cumulative CPU: 8.73 sec    HDFS
Read: 41000641 HDFS Write: 4521 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 730 msec
OK
AA      94.62
AAI     57.88
AAN     35.21
AAP     83.65
AAR     25.25
AAV     24.78
AB      94.94
```

```
.
.
.
.
.
.
.
```

```
Time taken: 29.413 seconds, Fetched: 203 row(s)
```

QUESTION 2

```
hive (navnath5028)>
create table customers(
cust_id INT,
firstname STRING,
lastname STRING,
age INT,
profession STRING
)
row format delimited
fields terminated by ','
stored as textfile;
OK
Time taken: 0.165 seconds
```

```
hive (navnath5028)> describe customer;
OK
cust_id          int
firstname        string
lastname         string
age              int
profession       string
```

```
hive (navnath5028)> show tables;
OK
airlines
airport
customer
```

```
hive (navnath5028)> select count(*) from customer;
Hadoop job information for Stage-1: number of mappers: 2; number of
reducers: 1
2022-06-20 10:08:06,185 Stage-1 map = 0%,   reduce = 0%
2022-06-20 10:08:17,707 Stage-1 map = 50%,   reduce = 0%, Cumulative
CPU 3.19 sec
2022-06-20 10:08:18,753 Stage-1 map = 100%,   reduce = 0%, Cumulative
CPU 6.11 sec
2022-06-20 10:08:34,344 Stage-1 map = 100%,   reduce = 100%,
Cumulative CPU 8.62 sec
MapReduce Total cumulative CPU time: 8 seconds 620 msec
Ended Job = job_1654490426372_5831
MapReduce Jobs Launched:
```

```
Stage-Stage-1: Map: 2   Reduce: 1   Cumulative CPU: 8.62 sec   HDFS
Read: 404450 HDFS Write: 105 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 620 msec
OK
10000
```

1) Write a program to find the count of customers for each profession

```
hive (navnath5028)>
```

select profession, count(profession) from customer group by profession;

```
Total MapReduce CPU Time Spent: 9 seconds 450 msec
```

```
OK
```

Accountant	199
Actor	202
Agricultural and food scientist	195
Architect	203
Artist	175
Athlete	196
Automotive mechanic	193
Carpenter	181
Chemist	209
Childcare worker	207
Civil engineer	193
Coach	201
Computer hardware engineer	204
Computer software engineer	216
Computer support specialist	222
Dancer	185
Designer	205
Doctor	197
Economist	189
Electrical engineer	192
Electrician	194
Engineering technician	204
Environmental scientist	176
Farmer	201
Financial analyst	198

2) Write a program to find the top 10 products sales wise

```
hive (navnath5028)>
create table sales(
txn_id INT,
txn_date STRING,
cust_id INT,
amount DOUBLE,
category STRING,
product STRING,
city STRING,
state STRING,
spendby STRING)
row format delimited
fields terminated by ','
stored as textfile;
OK
Time taken: 0.103 seconds
```

```
hive (navnath5028)> describe sales;
OK
txn_id          int
txn_date        string
cust_id         int
amount          double
category        string
product         string
city            string
state           string
spendby         string
Time taken: 0.062 seconds, Fetched: 9 row(s)
```

```
select product,round(sum(amount),2) as amt from sales group by product
order by amt desc limit 10;
```

```
Total MapReduce CPU Time Spent: 11 seconds 960 msec
OK
Yoga & Pilates 47804.94
```

```
Swing Sets          47204.14
Lawn Games          46828.44
Golf                46577.68
Cardio Machine Accessories  46485.54
Exercise Balls     45143.84
Weightlifting Belts  45111.68
Mahjong            44995.2
Basketball          44954.68
Beach Volleyball    44890.67
Time taken: 56.32 seconds, Fetched: 10 row(s)
```

3) Write a program to create partitioned table on category

```
hive (navnath5028)>
create table txnrecsByCat(
txn_id INT,
txn_date STRING,
cust_id INT,
amount DOUBLE,
product STRING,
city STRING,
state STRING,
spendby STRING)
partitioned by (category STRING)
row format delimited
fields terminated by ','
stored as textfile;
```

```
OK
Time taken: 0.165 seconds
```

QUESTION 3 PySpark

1) What was the highest number of people who traveled in which year?

```
dataRDD=sc.textFile("/user/bigdatamind43836/airlines.csv")
dataRDD2=dataRDD.map(lambda a : a.encode("ascii","ignore"))
header=dataRDD2.first()
```

```

dataRDD3=dataRDD2.map(lambda a : a != header)
dataRDD4=dataRDD3.map(lambda a : a.split(", "))
keyword=dataRDD4.map(lambda a : (a[0],int(a[3])))
count=keyword.reduceByKey(lambda a,b : a+b)
sortByval=count.sortBy(lambda a : -a[1])

for line in sortByval.collect():
    print(line)
...     print(i)
... output

('2007', 176299)

```

2) Identifying the highest revenue generation for which year

```

kvrdd2 = arrayrdd.map(lambda a : (a[0]+"
"+a[1],float(a[2])*int(a[3])))
counts2 = kvrdd2.reduceByKey(lambda a,b : a+b)
sort2 = counts2.sortBy(lambda a : -a[1])
>>> for i in sort2.take(1):
...     print(i)
...
('2014 4', 18819408.48)

```

3) Identifying the highest revenue generation for which year and quarter (Common group)

```

kvrdd = arrayrdd.map(lambda a : (a[0],float(a[2])*int(a[3])))
counts = kvrdd.reduceByKey(lambda a,b : a+b)
>>> sort = counts.sortBy(lambda a : -a[1])
>>> for i in sort.take(1):
...     print(i)
('2013', 66363208.71)

```


npbdh.cloudloka.com:4200

NYSE	AEA	2009-10-30	5.02	5.1	Mute	Stop Video	Participants	Chat	New Share	Pause Share	Annotate	More
NYSE	AEA	2009-10-29	5.18	5.49	5.01	5.07	439400	4.97				
NYSE	AEA	2009-10-28	4.91	5.14	4.77	5.02	267900	4.85				
NYSE	AEA	2009-10-27	5.06	5.16	4.82	4.9	353600	5.0				
NYSE	AEA	2009-10-26	5.13	5.36	4.95	5.05	266500	5.02				
NYSE	AEA	2009-10-23	5.34	5.48	5.05	5.07	199200	5.28				
NYSE	AEA	2009-10-22	5.26	5.37	5.17	5.33	434800	5.21				
NYSE	AEA	2009-10-21	5.3	5.61	5.25	5.26	508700	5.25				
NYSE	AEA	2009-10-20	5.75	5.76	5.24	5.3	365000	5.69				
NYSE	AEA	2009-10-19	5.77	5.98	5.62	5.75	489500	5.63				
NYSE	AEA	2009-10-16	5.98	5.98	5.6	5.69	229000	5.96				
NYSE	AEA	2009-10-15	6.07	6.15	6.02	6.02	247700	6.11				
NYSE	AEA	2009-10-14	6.1	6.19	6.05	6.17	356900	5.92				
NYSE	AEA	2009-10-13	6.15	6.15	5.73	5.98	181500	6.1				
NYSE	AEA	2009-10-12	6.1	6.22	6.08	6.16	186800	6.05				
NYSE	AEA	2009-10-09	5.98	6.12	5.8	6.11	301300	5.94				
NYSE	AEA	2009-10-08	6.01	6.08	5.97	6.0	314300	5.9				
NYSE	AEA	2009-10-07	5.91	5.96	5.69	5.96	276800	5.9				
NYSE	AEA	2009-10-06	5.88	5.96	5.81	5.96	444300	5.75				
NYSE	AEA	2009-10-05	5.57	5.93	5.51	5.81	279200	5.45				
NYSE	AEA	2009-10-02	5.4	5.55	5.3	5.51	351800	5.44				
NYSE	AEA	2009-10-01	5.56	5.67	5.36	5.5	330800	5.54				
NYSE	AEA	2009-09-30	5.44	5.63	5.21	5.6	248300	5.37				
NYSE	AEA	2009-09-29	5.54	5.56	5.38	5.42	277400	5.49				
NYSE	AEA	2009-09-28	5.2	5.65	5.12	5.55	399000	5.12				
NYSE	AEA	2009-09-25	5.49	5.49	5.12	5.17	285300	5.44				
NYSE	AEA	2009-09-24	5.8	5.87	5.41	5.5	230300	5.75				
NYSE	AEA	2009-09-23	6.15	6.17	5.81	5.81	238400	6.06				
NYSE	AEA	2009-09-22	5.91	6.26	5.91	6.12	186900	5.77				
NYSE	AEA	2009-09-21	5.8	5.94	5.7	5.83	318600	5.85				
NYSE	AEA	2009-09-18	5.93	5.93	5.74	5.91	195700	5.83				
NYSE	AEA	2009-09-17	5.77	5.92	5.75	5.89	240100	5.74				
NYSE	AEA	2009-09-16	5.8	5.82	5.66	5.8						

Time taken: 0.672 seconds, Fetched: 100 row(s)
hive (navnath5028)>

txns1.txt | custs.txt | Show all

Type here to search

eclipse-workspace - NYSE/src/AllTimeHigh.java - Eclipse IDE

File Edit Source Refactor Navigate Search Project Run Window Help

Package Explorer | StockVolume | AllTimeHigh.java

```
1 import java.io.*;
2 import org.apache.hadoop.io.Text;
3 import org.apache.hadoop.io.LongWritable;
4 import org.apache.hadoop.io.DoubleWritable;
5 import org.apache.hadoop.mapreduce.Job;
6 import org.apache.hadoop.mapreduce.Mapper;
7 import org.apache.hadoop.mapreduce.Reducer;
8 import org.apache.hadoop.conf.*;
9 import org.apache.hadoop.fs.*;
10 import org.apache.hadoop.mapreduce.lib.input.*;
11 import org.apache.hadoop.mapreduce.lib.output.*;
12
13
14 public class AllTimeHigh {
15
16     public static class MapClass extends Mapper<LongWritable,Text,Text,DoubleWritable>
17     {
18         private Text stock_id = new Text();
19         private DoubleWritable High = new DoubleWritable();
20
21         public void map(LongWritable key, Text value, Context context)
22         {
23             try{
24                 String[] str = value.toString().split(",");
25                 double high = Double.parseDouble(str[4]);
26                 stock_id.set(str[1]);
27                 High.set(high);
28
29                 //context.write(new Text(stock_id),new LongWritable(1));
30                 context.write(stock_id, High);
31             }
32             catch(Exception e)
33             {}
34         }
35     }
36
37     public static void main(String[] args) throws Exception {
38         System.out.println("Start Message");
39     }
40 }
```

Problems | Javadoc | Declaration

0 errors, 2 warnings, 0 others

Description	Resource	Path	Location	Type
Writable	Smart Insert	66:41:2211		

Type here to search

30°C | ENG | 15:04 | 20-06-2022

eclipse-workspace - NYSE/src/AllTimeHigh.java - Eclipse IDE

File Edit Source Refactor Navigate Search Project Run Window Help

Package Explorer

- CDAC
- NYSE
 - src
 - IRE System Library [JavaSE-1.8]
 - (default package)
 - AllTimeHigh.java
 - AllTimeHighExam.java
 - AllTimeLow.java
 - ClosingAvg.java
 - Offence.java
 - StockVolume.java
 - Student.java
- Referenced Libraries

```
27 stock_id.set(stock_id, High.set(high));
28
29 //context.write(new Text(stk[1]),new LongWritable(yo1));
30 context.write(stock_id, High);
31
32 catch(Exception e)
33 {
34     System.out.println(e.getMessage());
35 }
36
37 }
38
39 }
40
41 public static class ReduceClass extends Reducer<Text,DoubleWritable,Text,DoubleWritable>
42 {
43     private DoubleWritable result = new DoubleWritable();
44
45     public void reduce(Text key, Iterable<DoubleWritable> values,Context context) throws IOException, Interrupte
46     {
47         double maxVal=0;
48         double temp_val=0;
49         for (DoubleWritable value : values) {
50             temp_val = value.get();
51             if (temp_val > maxVal) {
52                 maxVal = temp_val;
53             }
54             result.set(maxVal);
55         }
56         context.write(key, result);
57         //context.write(key, new LongWritable(sum));
58     }
59 }
60
61 public static void main(String[] args) throws Exception {
```

Problems @ Javadoc Declaration

0 errors, 2 warnings, 0 others

Description	Resource	Path	Location	Type
Writable	Smart Insert	66:41:2211		

Type here to search

30°C 15:04 20-06-2022

Hive Intro Big L MOI Big L FTP Nuve bigd Sub Hue bigd

Not secure npbdh.cloudloka.cc

You are accessing a non-optimized Hue, please

Query Search saved documents...

navnath5028

Tables (20) +

Filter...

- airlines
- airport
- customer
- employee
- employee_header
- nyse
- routes
- testing
- tran
- txn_bucket
- txn_orc
- txn_parquet
- txnrecords
- txnrecsbycat
- txnrecsbycat3
- txnrecsbycat4
- txnrecsbycat_static
- txnrecsbycat_static2
- txnrecsbycat_static3
- txnrecsbyprod

Databases > navnath5028 > nyse

Overview Sample (0) Details

PROPERTIES

Table

Managed and stored in location

Created by bigdatamind43838 on 27/05/2022 19:48 +05:30

STATS

Files 0 Rows 0 Total size 0 B

Data last updated on 27/05/2022 19:48 +05:30

SCHEMA

Filter...

Column (9)	Type	Description	Sample
exchange_name	string	Add a description...	
stock_id	string	Add a description...	
stk_date	date	Add a description...	
open	double	Add a description...	
high	double	Add a description...	
low	double	Add a description...	

txns1.txt

custs.txt

Show all

Type here to search

30°C 15:06 20-06-2022

Hive x Intro x Big L x MOI x Big L x FTP x Nuve x bigd x Sub: x Hue x bigd x +

npbdh.cloudloka.cc

You are accessing a non-optimized Hue, please

Mute Stop Video Participants Chat New Share Pause Share Annotate More

You are screen sharing

Stop Share

Could not connect to ip: 10.1.1.204 on south-1.compute.int...
email:10000

AnalysisException: Un...
8.nyse.stk_date'.

28_Navnath Bhoskar_D8DA

File Browser

Back Home

Refresh

View as binary

Download

Last modified
27/05/2022 19:55 +05:30

User
bigdatamind43838

Group
hive

Size
39.09 MB

Mode
100644

/ user / hive / warehouse / navnath5028.db / nyse / NYSE.csv

```
NYSE, AEA, 2010-02-08, 4.42, 4.42, 4.21, 4.24, 205500, 4.24
NYSE, AEA, 2010-02-05, 4.42, 4.54, 4.22, 4.41, 194300, 4.41
NYSE, AEA, 2010-02-04, 4.55, 4.69, 4.39, 4.42, 233800, 4.42
NYSE, AEA, 2010-02-03, 4.65, 4.69, 4.50, 4.55, 182100, 4.55
NYSE, AEA, 2010-02-02, 4.74, 5.00, 4.62, 4.66, 222700, 4.66
NYSE, AEA, 2010-02-01, 4.84, 4.92, 4.68, 4.75, 194800, 4.75
NYSE, AEA, 2010-01-29, 4.97, 5.05, 4.76, 4.83, 222900, 4.83
NYSE, AEA, 2010-01-28, 5.12, 5.22, 4.81, 4.98, 283100, 4.98
NYSE, AEA, 2010-01-27, 4.82, 5.16, 4.79, 5.09, 243500, 5.09
NYSE, AEA, 2010-01-26, 5.18, 5.18, 4.81, 4.84, 554800, 4.84
NYSE, AEA, 2010-01-25, 5.42, 5.48, 5.20, 5.22, 257300, 5.22
NYSE, AEA, 2010-01-22, 5.52, 5.59, 5.31, 5.37, 260800, 5.37
NYSE, AEA, 2010-01-21, 5.67, 5.74, 5.37, 5.51, 264300, 5.51
NYSE, AEA, 2010-01-20, 5.65, 5.70, 5.53, 5.66, 244600, 5.66
NYSE, AEA, 2010-01-19, 5.54, 5.70, 5.54, 5.69, 368000, 5.69
NYSE, AEA, 2010-01-15, 5.48, 5.55, 5.33, 5.54, 435500, 5.54
NYSE, AEA, 2010-01-14, 5.41, 5.50, 5.39, 5.41, 272200, 5.41
NYSE, AEA, 2010-01-13, 5.50, 5.50, 5.41, 5.45, 176400, 5.45
NYSE, AEA, 2010-01-12, 5.47, 5.51, 5.41, 5.46, 233100, 5.46
NYSE, AEA, 2010-01-11, 5.64, 5.64, 5.49, 5.55, 178900, 5.55
```

txns1.txt

custs.txt

Show all

Type here to search

Hive x Intro x Big L x MOI x Big L x FTP x Nuve x bigd x Sub: x Hue x bigd x +

npbdh.cloudloka.com:4200

Mute Stop Video Participants Chat New Share Pause Share Annotate More

You are screen sharing

Stop Share

```
> txn_date STRING,
>
> cust_id INT,
>
> amount DOUBLE,
>
> category STRING,
>
> product STRING,
>
> city STRING,
>
> state STRING,
>
> spendby STRING)
>
> row format delimited
>
> fields terminated by ','
>
> stored as textfile;
```

OK

Time taken: 0.165 seconds

hive (navnath5028)>

txns1.txt

custs.txt

Show all

Type here to search

29°C

15:33

20-06-2022

2022-06-20 16:15:50,523 Stage-1 Map: 100%

MapReduce Total cumulative CPU time: 9 seconds 450 msec

Ended Job = job_1654490426372_5861

MapReduce Jobs Launched:

Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 9.45 sec HDFS Read: 405218 HDFS Write: 1584 HDFS EC Read: 0 SUCCESS

Total MapReduce CPU Time Spent: 9 seconds 450 msec

OK

Accountant	199
Actor	202
Agricultural and food scientist	195
Architect	203
Artist	175
Athlete	196
Automotive mechanic	193
Carpenter	181
Chemist	209
Childcare worker	207
Civil engineer	193
Coach	201
Computer hardware engineer	204
Computer software engineer	216
Computer support specialist	222
Dancer	185
Designer	205
Doctor	197
Economist	189
Electrical engineer	192
Electrician	194
Engineering technician	204
Environmental scientist	176
Farmer	201
Financial analyst	198
Firefighter	217
Human resources assistant	212
Judge	196
Lavender	212

txns1.txt | custs.txt

Type here to search

29°C 15:46 20-06-2022

> product STRING,

>

> city STRING,

>

> state STRING,

>

> spendby STRING)

>

> row format delimited

>

> fields terminated by ','

>

> stored as textfile;

OK

Time taken: 0.103 seconds

hive (navnath5028)> describe sales;

OK

txn_id	int
txn_date	string
cust_id	int
amount	double
category	string
product	string
city	string
state	string
spendby	string

Time taken: 0.062 seconds, Fetched: 9 row(s)

hive (navnath5028)>

txns1.txt | custs.txt

Type here to search

29°C 15:54 20-06-2022

```
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
22/06/20 11:09:28 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
22/06/20 11:09:28 INFO client.RMProxy: Connecting to ResourceManager at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
Starting Job = job_1654490426372_5980, Tracking URL = http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/application_1654490426372_5980/
Kill Command = /opt/cloudera/parcels/CDH-6.2.1-1.cdh6.2.1.p0.1425774/lib/hadoop/bin/hadoop job -kill job_1654490426372_5980
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2022-06-20 11:09:37,496 Stage-2 map = 0%, reduce = 0%
2022-06-20 11:09:45,732 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.98 sec
2022-06-20 11:09:53,927 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 1.98 sec
MapReduce Total cumulative CPU time: 4 seconds 580 msec
Ended Job = job_1654490426372_5980
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.38 sec HDFS Read: 4426996 HDFS Write: 4865 HDFS EC Read: 0 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 4.58 sec HDFS Read: 10386 HDFS Write: 436 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 960 msec
OK
Yoga & Pilates 47804.94
Swing Sets 47204.14
Lawn Games 46828.44
Golf 46577.68
Cardio Machine Accessories 46485.54
Exercise Balls 45143.84
Weightlifting Belts 45111.68
Mahjong 44995.2
Basketball 44954.68
Beach Volleyball 44890.67
Time taken: 56.32 seconds, Fetched: 10 row(s)
hive (navnath5028)>
```



```
at org.apache.hadoop.util.RunJar.run(RunJar.java:227)
at org.apache.hadoop.util.RunJar.main(RunJar.java:227)
FAILED: ParseException line 1:0 cannot recognize input near 'create' 'table' 'txnrecsByCat'
hive (navnath5028)> create table txnrecsByCat(
>
>   > txn_id INT,
>
>   > txn_date STRING,
>
>   > cust_id INT,
>
>   > amount DOUBLE,
>
>   > product STRING,
>
>   > city STRING,
>
>   > state STRING,
>
>   > spendby STRING)
>
>   > partitioned by (category STRING)
>
>   > row format delimited
>
```

pgAdmin 4

File Object Tools Help

Browser SQL Properties nava_services/postgres@PostgreSQL 14

Servers (1)

- PostgreSQL 14
 - Databases (9)
 - assignments
 - cdac_2022
 - exam
 - imdb_database
 - ipl
 - joins
 - match_day
 - nava_services
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wra
 - Languages
 - Publications
 - Schemas
 - Subscriptions
 - postgres
 - Login/Group Roles
 - Tablespaces

SQL

```
nava_services=# UPDATE quant_sip set nav_date = '30-06-2021' WHERE nav_date = '30-06-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '27-07-2021' WHERE nav_date = '27-07-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '24-09-2021' WHERE nav_date = '24-09-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '20-10-2021' WHERE nav_date = '20-10-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '25-10-2021' WHERE nav_date = '25-10-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '17-11-2021' WHERE nav_date = '17-11-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '24-11-2021' WHERE nav_date = '24-11-2020';
UPDATE 1
nava_services=# UPDATE quant_sip set nav_date = '27-12-2021' WHERE nav_date = '27-12-2020';
UPDATE 1
nava_services=# select * from quant_sip;
 serial_id | amount | units | nav_price | nav_date
-----
1 | 500 | 3.73 | 134.0355 | 2020-12-02
2 | 500 | 3.359 | 148.8651 | 2021-01-07
3 | 500 | 3.186 | 156.9075 | 2021-02-12
4 | 500 | 3.055 | 163.6462 | 2021-03-08
5 | 500 | 2.867 | 174.3931 | 2021-04-22
6 | 500 | 2.481 | 201.5054 | 2020-06-04
7 | 500 | 2.432 | 205.6167 | 2021-06-30
8 | 500 | 2.303 | 217.1145 | 2021-07-27
9 | 500 | 2.195 | 227.7731 | 2021-09-24
10 | 500 | 2.099 | 238.2115 | 2021-10-20
11 | 500 | 2.151 | 232.4872 | 2021-10-25
12 | 500 | 2.066 | 242.0053 | 2021-11-17
13 | 500 | 2.089 | 239.3373 | 2021-11-24
14 | 500 | 2.143 | 233.2965 | 2021-12-27
(14 rows)
```

nava_services=# SE

npbdh.cloudloka.com:4200

Coach 201
Computer hardware engineer 204
Computer software engineer 216
Computer support specialist 222
Dancer 185
Designer 205
Doctor 197
Economist 189
Electrical engineer 192
Electrician 194
Engineering technician 204
Environmental scientist 176
Farmer 201
Financial analyst 198
Firefighter 217
Human resources assistant 212
Judge 196
Lawyer 212
Librarian 218
Loan officer 221
Musician 205
Nurse 192
Pharmacist 213
Photographer 222
Physicist 201
Pilot 212
Police officer 210
Politician 228
Psychologist 194
Real estate agent 191
Recreation and fitness worker 210
Reporter 200
Secretary 200
Social Worker 1
Social worker 212

txns1.txt

custs.txt

Show all

```
txnrrecsbycat3
txnrrecsbycat4
txnrrecsbycat_static
txnrrecsbycat_static2
txnrrecsbycat_static3
txnrrecsbyprod
Time taken: 0.112 seconds, Fetched: 22 row(s)
hive (navnath5028)> describe txnrrecsbycat;
OK
txnrno                int
txnrdate              string
custno                int
amount                double
product               string
city                  string
state                  string
spendby               string
category              string

# Partition Information
# col_name             data_type          comment
category               string
Time taken: 0.154 seconds, Fetched: 14 row(s)
hive (navnath5028)> create table txnrrecsByCat(
>   txn_id INT,
>   txn_date STRING,
>   cust_id INT,
>   amount DOUBLE,
>   product STRING,
>   city STRING,
>   state STRING,
>   spendby STRING)
> partitioned by (category STRING)
> row format delimited
```